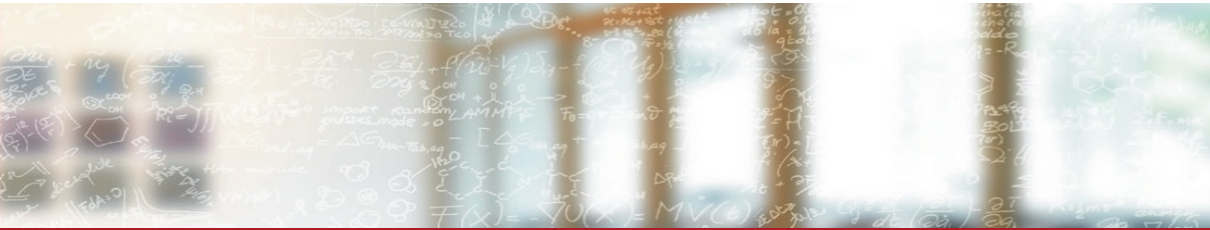




CSCS

Centro Svizzero di Calcolo Scientifico
Swiss National Supercomputing Centre

ETH zürich



Introduction to the Piz Daint environment

Directive Based GPU Programming

Vasileios Karakasis, CSCS

May 14–15, 2018

Overview

- Accessing CSCS
- Compiling my code
- Running my code
- Editing my code
- Transferring files from/to CSCS
- Repository of the course

Piz Daint

Computing nodes

Piz Daint is a hybrid cluster of Cray XC40/XC50 nodes

- Hybrid nodes (XC50)
 - 5320 total
 - Intel® Xeon® E5-2690 v3 @ 2.60GHz (12 cores, 64GB RAM, Haswell)
 - NVIDIA® Tesla® P100 16GB (Pascal)
- Multicore nodes
 - 1813 nodes
 - Two Intel® Xeon® E5-2695 v4 @ 2.10GHz (2 x 18 cores, 64/128 GB RAM, Broadwell)
- Login nodes
 - 5 total
 - Intel® Xeon® CPU E5-2650 v3 @ 2.30GHz (10 cores, 256 GB RAM, Haswell)
- Aries routing and communications ASIC, and Dragonfly network topology

Piz Daint

Filesystems

- `/scratch`: High performance Lustre filesystem accessible from the computing nodes
 - Environment variable `$SCRATCH` points to it
 - Total capacity: 6.2 PB
 - Must be used for heavy I/O
- `/users`: GPFS filesystem for the users' homes
- `/project`, `/store`: Long-term storage for computational projects

More on https://user.cscs.ch/storage/file_systems/

Accessing Piz Daint

- Accessible through SSH
- Piz Daint is not directly accessible from the outside world:
 - `ela` → `daint10x` → `nidxxxxxx`

Two-steps process:

1. Login to the frontend, forwarding X11 (will be needed the second day)
2. Move to the login nodes of Piz Daint

```
# Login to the frontend first  
ssh -Y courseXX@ela.cscs.ch  
ssh -Y daint
```

Programming Environments

Cray Linux Programming Environment

- 4 compilers available: CCE, GNU, INTEL, PGI
- 4 predefined Programming Environments:
 - PrgEnv-cray (default), PrgEnv-gnu, PrgEnv-intel, PrgEnv-pgi
 - `echo $PE_ENV` to get the current PrgEnv
- 3 wrappers available: `ftn` (Fortran), `cc` (C), `CC` (C++)
 - Required for compiling MPI programs
 - They set appropriate optimisation flags for the target architecture (CPU or GPU)
 - They provide a sort of portability across the programming environments

Managing programming environments

Daint uses *Environment Modules* (TMod) for managing the programming environments and the software packages:

- Dynamic modification of a user's environment via *modulefiles*.
- All programming environments and software on Daint is available through modules.
- The compiler wrappers will detect the loaded programming environment and automatically set the correct flags and libraries.

Managing programming environments

Listing modules

– `module list`

```
2. ssh
course00@daint104:~> module list
Currently Loaded Modulefiles:
  1) modules/3.2.10.6
  2) eproxy/2.0.16-6.0.4.1_3.1__g001b199.ari
  3) cce/8.6.1
  4) craype-network-aries
  5) craype/2.5.12
  6) cray-libsci/17.06.1
  7) udreg/2.3.2-6.0.4.0_12.2__g2f9c3ee.ari
  8) ugni/6.0.14-6.0.4.0_14.1__ge7db4a2.ari
  9) pmi/5.0.12
 10) dmapp/7.1.1-6.0.4.0_46.2__gb8abda2.ari
 11) gni-headers/5.0.11-6.0.4.0_7.2__g7136988.ari
 12) xpmem/2.2.2-6.0.4.0_3.1__g43b0535.ari
 13) job/2.2.2-6.0.4.0_8.2__g3c644b5.ari
 14) dvs/2.7_2.2.32-6.0.4.1_7.1__ged1923a
 15) alps/6.4.1-6.0.4.0_7.2__g86d0f3d.ari
 16) rca/2.2.11-6.0.4.0_13.2__g84de67a.ari
 17) atp/2.1.1
 18) perftools-base/6.5.1
 19) PrgEnv-cray/6.0.4
 20) cray-mpich/7.6.0
 21) slurm/17.02.9+git20180119.b04278-1
 22) craype-haswell
 23) xalt/daint-2016.11
course00@daint104:~> █
```


Managing programming environments

Switching programming environments

- Switch to the PGI programming environment
- module switch

```
2. ssh
course00@daint104:~> module switch PrgEnv-cray/6.0.4 PrgEnv-pgi
course00@daint104:~> module list
Currently Loaded Modulefiles:
 1) modules/3.2.10.6                  12) pmi/5.0.12
 2) eproxy/2.0.16-6.0.4.1_3.1__g001b199.ari 13) dmapp/7.1.1-6.0.4.0_46.2__gb8abda2.ari
 3) pgi/17.5.0                       14) gni-headers/5.0.11-6.0.4.0_7.2__g7136988.ari
 4) craype-haswell                  15) xpmem/2.2.2-6.0.4.0_3.1__g43b0535.ari
 5) craype-network-aries            16) job/2.2.2-6.0.4.0_8.2__g3c644b5.ari
 6) craype/2.5.12                   17) dvs/2.7_2.2.32-6.0.4.1_7.1__ged1923a
 7) cray-mpich/7.6.0                18) alps/6.4.1-6.0.4.0_7.2__g86d0f3d.ari
 8) slurm/17.02.9+git20180119.b04278-1 19) rca/2.2.11-6.0.4.0_13.2__g84de67a.ari
 9) xalt/daint-2016.11              20) atp/2.1.1
10) udreg/2.3.2-6.0.4.0_12.2__g2f9c3ee.ari 21) perftools-base/6.5.1
11) ugni/6.0.14-6.0.4.0_14.1__ge7db4a2.ari 22) PrgEnv-pgi/6.0.4
course00@daint104:~> ftn -V

pgf90 17.5-0 64-bit target on x86-64 Linux -tp haswell-64
PGI Compilers and Tools
Copyright (c) 2017, NVIDIA CORPORATION. All rights reserved.
course00@daint104:~> █
```

Managing programming environments

Switching back to the Cray programming environment

```
2. ssh
course00@daint104:~> module switch PrgEnv-pgi/6.0.4 PrgEnv-cray
course00@daint104:~> module list
Currently Loaded Modulefiles:
 1) modules/3.2.10.6
 2) eproxy/2.0.16-6.0.4.1_3.1__g001b199.ari
 3) slurm/17.02.9+git20180119.b04278-1
 4) xalt/daint-2016.11
 5) cce/8.6.1
 6) craype-haswell
 7) craype-network-aries
 8) craype/2.5.12
 9) cray-mpich/7.6.0
10) cray-libsci/17.06.1
11) udreg/2.3.2-6.0.4.0_12.2__g2f9c3ee.ari
12) ugni/6.0.14-6.0.4.0_14.1__ge7db4a2.ari
13) pmi/5.0.12
14) dmapp/7.1.1-6.0.4.0_46.2__gb8abda2.ari
15) gni-headers/5.0.11-6.0.4.0_7.2__g7136988.ari
16) xpmem/2.2.2-6.0.4.0_3.1__g43b0535.ari
17) job/2.2.2-6.0.4.0_8.2__g3c644b5.ari
18) dvs/2.7.2.2.32-6.0.4.1_7.1__ged1923a
19) alps/6.4.1-6.0.4.0_7.2__g86d0f3d.ari
20) rca/2.2.11-6.0.4.0_13.2__g84de67a.ari
21) atp/2.1.1
22) perftools-base/6.5.1
23) PrgEnv-cray/6.0.4
course00@daint104:~> ftn -V
Cray Fortran : Version 8.6.1 Sun May 13, 2018 19:16:04
course00@daint104:~> █
```

Managing programming environments

Loading and unloading modules

- `module load [MODULE_NAME]`
- `module unload [MODULE_NAME]`

```
2. ssh
course00@daint104:~> module load cray-hdf5
course00@daint104:~> which h5dump
/opt/cray/pe/hdf5/1.10.0.3/bin/h5dump
course00@daint104:~> module unload cray-hdf5
course00@daint104:~> which h5dump
which: no h5dump in (/opt/cray/pe/perftools/6.5.1/bin:/opt/cray/pe/papi/5.5.1.2/bin:/opt/cray/rca/2.2.11-6.0.4.0_13.2__g84de67a.ari/bin:/opt/cray/alps/6.4.1-6.0.4.0_7.2__g86d0f3d.ari/sbin:/opt/cray/job/2.2.2-6.0.4.0_8.2__g3c644b5.ari/bin:/opt/cray/pe/mpt/7.6.0/gni/bin:/opt/cray/pe/craype/2.5.12/bin:/opt/cray/pe/cce/8.6.1/binutils/x86_64/x86_64-pc-linux-gnu/bin:/opt/cray/pe/cce/8.6.1/binutils/cross/x86_64-aarch64/aarch64-unknown-linux-gnu/..bin:/opt/cray/pe/cce/8.6.1/utlis/x86_64/bin:/apps/daint/UES/xalt/0.7.6/bin:/opt/slurm/17.02.9+git20180119.b04278/bin:/opt/cray/elogin/eproxy/2.0.16-6.0.4.1_3.1__g001b199.ari/bin:/opt/cray/pe/modules/3.2.10.6/bin:/opt/slurm/default/bin:/apps/daint/system/bin:/apps/common/system/bin:/users/course00/bin:/usr/local/bin:/usr/bin:/bin:/usr/bin/X11:/usr/lib/mit/bin:/usr/lib/mit/sbin:/opt/cray/pe/bin)
course00@daint104:~> h5dump
If 'h5dump' is not a typo you can use command-not-found to lookup the package that contains it, like this:
  cnf h5dump
course00@daint104:~> █
```

Managing programming environments

Checking available modules

- The `daint-gpu` makes available the CSCS software stack built for the hybrid nodes of the system

```
2. ssh
course00@daint104:~> module load daint-gpu
course00@daint104:~> module avail

----- /apps/daint/UES/jenkins/6.0.UP04/gpu/easybuild/modules/all -----
Amber/16-CrayGNU-17.08-cuda-8.0(default)      Score-P/3.1-CrayIntel-17.08
Amber/16-CrayGNU-17.08-parallel              Score-P/3.1-CrayPGI-17.08
Amber/16-CrayGNU-17.08-serial                Spark/1.6.0(default)
Boost/1.65.0-CrayGNU-17.08                  TensorFlow/1.2.1-CrayGNU-17.08-cuda-8.0-python3(default)
Boost/1.65.0-CrayGNU-17.08-python2          TensorFlow/1.3.0-CrayGNU-17.08-cuda-8.0-python3
Boost/1.65.0-CrayGNU-17.08-python3(default) TensorFlow/1.4.1-CrayGNU-17.08-cuda-8.0-python3
CD0/1.9.0-CrayGNU-17.08(default)            TensorFlow/1.4.1-CrayGNU-17.12-cuda-8.0-python3
CD0/1.9.0-CrayIntel-17.08                  TensorFlow/1.7.0-CrayGNU-17.12-cuda-8.0-python3
CMake/3.9.1                                Theano/0.9.0-CrayGNU-17.08-cuda-8.0-python2
CMake/3.10.1                               Theano/0.9.0-CrayGNU-17.08-cuda-8.0-python3(default)
CP2K/5.0r18043-CrayGNU-17.08-cuda-8.0      Theano/1.0.1-CrayGNU-17.08-python2
CP2K/5.1-CrayGNU-17.08-cuda-8.0(default)    Theano/1.0.1-CrayGNU-17.08-python3
CPMD/4.1-CrayIntel-17.08g(default)          VASP/5.4.4-CrayIntel-17.08-cuda-8.0(default)
Charm++/6.8.0-CrayIntel-17.08(default)      VMD/1.9.3-egl
Circos/0.69-6-Perl-5.26.1-bare(default)    VMD/1.9.3-ogl
Dimemas/5.3.3(default)                    VTK/8.0.1-CrayGNU-17.08-python2
EasyBuild-custom/cscs                     VTK/8.0.1-CrayGNU-17.08-python3
Extrac/3.5.1-CrayGNU-17.08                 VTK/8.0.1-EGL-CrayGNU-17.08-python2
Extrac/3.5.1-CrayIntel-17.08               VTK/8.0.1-EGL-CrayGNU-17.08-python3
```

Managing programming environments

Checking available modules

- Check available versions of a software




```
course00@daint104:~> module avail gcc

----- /opt/modulefiles -----
gcc/4.9.3      gcc/5.3.0(default) gcc/6.1.0      gcc/6.2.0      gcc/7.1.0      gcc/7.2.0
course00@daint104:~> 
```

Managing programming environments

Get information about a module

- Environment variables set, paths etc.

A terminal window titled '2. ssh' showing the output of the 'module show gcc' command. The output lists various environment variables and their values for the gcc module, including paths for bin, man, info, and library, as well as GCC and GNU versions.

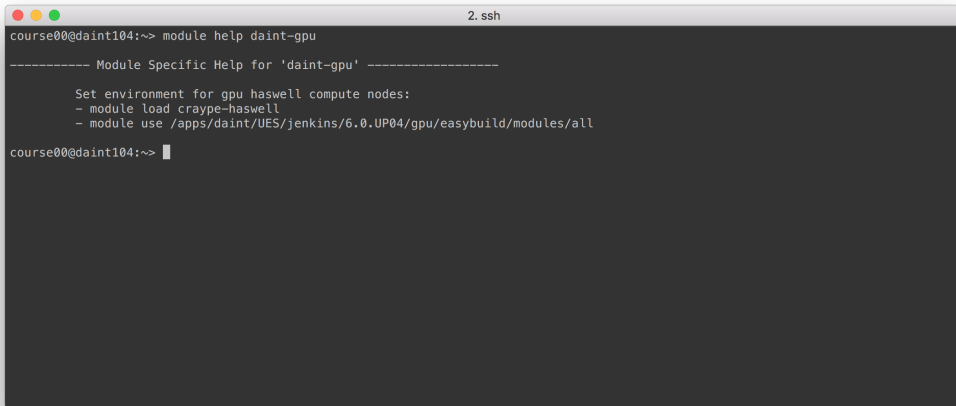
```
course00@daint104:~> module show gcc
-----
/opt/modulefiles/gcc/5.3.0:

conflict      gcc
conflict      gcc-cross-aarch64
prepend-path  PATH /opt/gcc/5.3.0/bin
prepend-path  MANPATH /opt/gcc/5.3.0/snos/share/man
prepend-path  INFOPATH /opt/gcc/5.3.0/snos/share/info
prepend-path  LD_LIBRARY_PATH /opt/gcc/5.3.0/snos/lib64
setenv        GCC_PATH /opt/gcc/5.3.0
setenv        GCC_VERSION 5.3.0
setenv        GNU_VERSION 5.3.0
-----

course00@daint104:~> █
```

Managing programming environments

Get help for a module



A terminal window titled "2. ssh" showing the command `module help daint-gpu` and its output. The output displays the module-specific help for 'daint-gpu', which includes instructions to set the environment for GPU Haswell compute nodes by loading the 'craype-haswell' module and using the specific module path for the UES/Jenkins environment.

```
course00@daint104:~> module help daint-gpu

----- Module Specific Help for 'daint-gpu' -----

Set environment for gpu haswell compute nodes:
- module load craype-haswell
- module use /apps/daint/UES/jenkins/6.0.UP04/gpu/easybuild/modules/all

course00@daint104:~> █
```

Running on Piz Daint

The job scheduler

Piz Daint uses native SLURM for running jobs on the compute nodes. There are two ways of submitting a job:

1. Interactively from the login nodes using the `srun` command.
2. By submitting a job script using the `sbatch` command.

Running on Piz Daint

Using the `srun` command

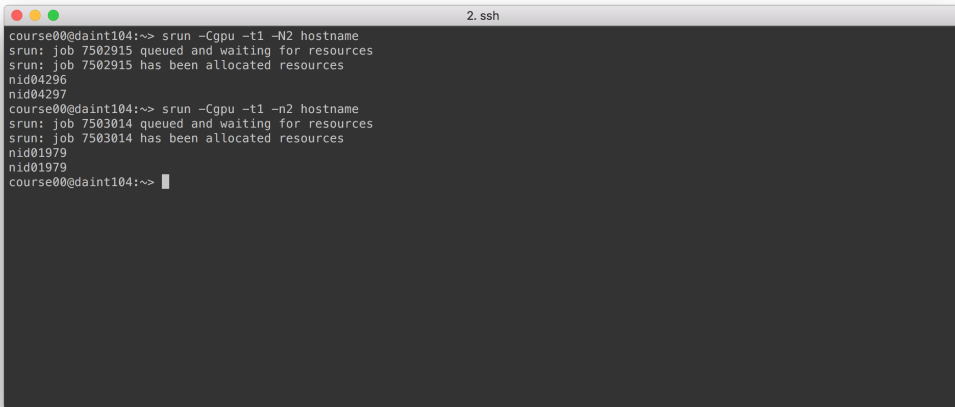
Necessary and useful options:

- `-C gpu`: requests allocation on the hybrid (GPU) nodes (required)
- `--reservation=openacc`: the reservation for our course to avoid waiting times
 - Reservation is valid for the two days of the course from 07:00-19:00.
- `-N 2`: number of compute nodes (default is 1)
- `-n 2`: number of MPI tasks (default is 1)
- `-t 5`: maximum duration of the job (default is 30min)
 - Allows to get an allocation quicker
 - Job will be killed if time limit is reached
 - Maximum time slot for a job is 24h

More on https://user.cscs.ch/getting_started/running_jobs/

Running on Piz Daint

Using the srun command

A terminal window titled "2. ssh" with a dark background and light text. It shows the execution of the srun command to run a hostname command on a GPU node. The output shows the job being queued, resources being allocated, and the node ID (nid04296 and nid04297) being printed. The prompt returns to the user after the command completes.

```
course00@daint104:~> srun -Cgpu -t1 -N2 hostname
srun: job 7502915 queued and waiting for resources
srun: job 7502915 has been allocated resources
nid04296
nid04297
course00@daint104:~> srun -Cgpu -t1 -n2 hostname
srun: job 7503014 queued and waiting for resources
srun: job 7503014 has been allocated resources
nid01979
nid01979
course00@daint104:~> █
```

Running on Piz Daint

Using the sbatch command

```
2. ssh
course00@daint104:~> cat job.sh
#!/bin/bash
#SBATCH -J 'my_first_job'
#SBATCH -C gpu
#SBATCH -N 2
#SBATCH -t 1
#SBATCH -o myjob.out
#SBATCH -e myjob.err

echo "My job id is $SLURM_JOB_ID"
hostname
course00@daint104:~> sbatch job.sh
Submitted batch job 7503300
course00@daint104:~> squeue -j 7503300
  JOBID  USER  ACCOUNT      NAME  ST REASON   START_TIME          TIME  TIME_LEFT  NODES  CPUS
  7503300 course00  crs03  my_first_job  CG None    20:44:18          0:06       0:54       2    48
course00@daint104:~> squeue -u $USER
  JOBID  USER  ACCOUNT      NAME  ST REASON   START_TIME          TIME  TIME_LEFT  NODES  CPUS
course00@daint104:~> █
```

Running on Piz Daint

Using the `sbatch` command – Examining the output

```
2. ssh
course00@daint104:~> cat myjob.err
course00@daint104:~> cat myjob.out
My job id is 7503300
nid03663

Batch Job Summary Report for Job "my_first_job" (7503300) on daint
-----
Submit          Eligible          Start          End          Elapsed  Timelimit
-----
2018-05-13T20:44:17 2018-05-13T20:44:17 2018-05-13T20:44:18 2018-05-13T20:44:24 00:00:06 00:01:00
-----

Username  Account  Partition  NNodes  Energy
-----
course00  crs03    normal     2        40 joules
-----

Scratch File System  Files  Quota
-----
/scratch/snx3000     2     1000000
-----

course00@daint104:~> █
```

Running on Piz Daint

Other useful commands

- `squeue [OPTIONS]`: Check the status of the system job queue
 - Useful options: `-u [USERNAME]`, `-j [JOBID]`
- `scancel [JOBID]`: Cancel a job
- `scontrol`: Detailed information about partitions, reservations, computing nodes etc.

Running on Piz Daint

Other useful commands

```
2. ssh
course00@daint104:~> scontrol show reservation openacc
ReservationName=openacc StartTime=Tomorr 07:00 EndTime=Tue 19:00 Duration=1-12:00:00
Nodes=nid0[7104-7119] NodeCnt=16 CoreCnt=192 Features=(null) PartitionName=(null) Flags=SPEC_NODES
TRES=cpu=384
Users=(null) Accounts=root,csstaff,crs03 Licenses=(null) State=INACTIVE BurstBuffer=(null) Watts=n/a

course00@daint104:~> scontrol show partition normal
PartitionName=normal
  AllowGroups=ALL AllowAccounts=ALL AllowQos=ALL
  AllocNodes=ALL Default=YES QoS=N/A
  DefaultTime=00:30:00 DisableRootJobs=NO ExclusiveUser=NO GraceTime=0 Hidden=NO
  MaxNodes=2400 MaxTime=1-00:00:00 MinNodes=1 LLN=NO MaxCPUsPerNode=UNLIMITED
  Nodes=nid[00004-00007,00012-00024,00026-00062,00064-00067,00072-00126,00128-00190,00192-00195,00200-00254,00260-00318,00320-003
23,00328-00382,00388-00446,00456-00510,00516-00568,00571-00574,00576-00579,00584-00638,00644-00702,00704-00707,00712-00766,00772-0
0830,00832-00835,00840-00894,00900-00958,00960-00963,00968-01022,01028-01086,01088-01150,01152-01192,01194-01214,01216-01260,01262
-01284,01286-01791,01804-01823,01868-01887,01920-01935,01940-01967,01972-02319,02324-02351,02356-02703,02708-02735,02740-03087,030
92-03119,03124-03471,03476-03503,03512-03855,03860-03887,03892-04239,04244-04271,04280-07295]
  PriorityJobFactor=10 PriorityTier=20 RootOnly=NO ReqResv=NO OverSubscribe=EXCLUSIVE
  OverTimeLimit=NONE PreemptMode=OFF
  State=UP TotalCPUs=252408 TotalNodes=7047 SelectTypeParameters=NONE
  DefMemPerNode=UNLIMITED MaxMemPerNode=UNLIMITED

course00@daint104:~> |
```

Editing files

- vim or gvim (X version)
- emacs -nw or just emacs (X version)
- gedit (X only)

Moving data to/from CSCS

- **scp**: Remote copy over SSH
 - Getting a file: `scp course00@ela.cscs.ch:remotefile localfile`
 - Getting a directory: `scp -r course00@ela.cscs.ch:remotedir localdir`
 - Sending a file: `scp localfile course00@ela.cscs.ch:remotefile`
 - Sending a directory: `scp localdir course00@ela.cscs.ch:remotedir`
- **rsync**: Synchronize files remotely over SSH
 - `rsync -avz course00@ela.cscs.ch:remotedir/ localdir/`
 - `rsync -avz localdir/ course00@ela.cscs.ch:remotedir/`
 - Pay attention to the slashes! *rsync behaves differently with or without slashes.*